

Mini Project Report

Your Name

1 Library Imports and Dataset Download

The first step in this project involves importing the required libraries and downloading the dataset using the `gdown` module. Below is the Python code used:

```
1 # Install and download dataset from Google Drive
2 !pip install --upgrade --no-cache-dir gdown
3 !gdown 1A7NRguAV3PZdxK6zsDDSQa1sqwz9IfU7
4
5 # Data handling and visualization
6 import pandas as pd
7 import matplotlib.pyplot as plt
8 import seaborn as sns
9 import numpy as np
10
11 # Data preprocessing and splitting
12 from sklearn.model_selection import train_test_split
13 from sklearn.preprocessing import MinMaxScaler
14
15 # PyTorch modules for neural network
16 import torch
17 import torch.nn as nn
18 import torch.optim as optim
19
20 # Evaluation metrics
21 from sklearn.metrics import r2_score
22 from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
```

Listing 1: Downloading Dataset and Importing Libraries

Explanation

- `gdown` is used to download the dataset directly from Google Drive using its file ID.
- Libraries such as `pandas`, `matplotlib`, and `seaborn` are used for data manipulation and visualization.
- `sklearn` provides tools for scaling, splitting, and evaluating models.
- `torch` is used for building and training deep learning models.

Output

This section does not produce any visual output but sets up all necessary dependencies for further development.

2 Data Loading and Initial Exploration

The dataset used in this project contains records related to graduate admission. Below is the code used to load and inspect the dataset.

```

1 # Load dataset
2 df = pd.read_csv("Admission_Predict.csv")
3
4 # Display structure and data types
5 df.info()
6
7 # Display summary statistics
8 df.describe()
9
10 # Display first few rows
11 df.head()

```

Listing 2: Loading and Exploring the Dataset

Explanation

- `pd.read_csv(...)` loads the dataset into a `DataFrame`. `df.info()` provides metadata: column names, datatypes, non-null counts, and memory usage.
- `df.describe()` generates summary statistics (mean, std, min, max, etc.) for numerical columns.
- `df.head()` shows the first five entries of the dataset for an initial look at the data.

DataFrame Info Output

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 400 entries, 0 to 399
Data columns (total 9 columns):
 #   Column                Non-Null Count  Dtype
---  -
0   Serial No.            400 non-null   int64
1   GRE Score              400 non-null   int64
2   TOEFL Score            400 non-null   int64
3   University Rating      400 non-null   int64
4   SOP                    400 non-null   float64
5   LOR                    400 non-null   float64
6   CGPA                   400 non-null   float64
7   Research                400 non-null   int64
8   Chance of Admit        400 non-null   float64
dtypes: float64(4), int64(5)
memory usage: 28.3 KB

```

Sample Data (First 5 Rows)

| Serial No. | GRE | TOEFL | Univ Rating | SOP | LOR | CGPA | Res. | Admit Chance |
|------------|-----|-------|-------------|-----|-----|------|------|--------------|
| 1 | 337 | 118 | 4 | 4.5 | 4.5 | 9.65 | 1 | 0.92 |
| 2 | 324 | 107 | 4 | 4.0 | 4.5 | 8.87 | 1 | 0.76 |
| 3 | 316 | 104 | 3 | 3.0 | 3.5 | 8.00 | 1 | 0.72 |
| 4 | 322 | 110 | 3 | 3.5 | 2.5 | 8.67 | 1 | 0.80 |
| 5 | 314 | 103 | 2 | 2.0 | 3.0 | 8.21 | 0 | 0.65 |

Table 1: First 5 rows of the dataset

3 Feature Correlation Analysis

To understand the relationships between different features and the target variable (Chance of Admit), a correlation matrix was computed and visualized using a heatmap.

```

1 # Compute correlation matrix for numerical features
2 correlation_matrix = df.corr(numeric_only=True)
3
4 # Plotting the heatmap
5 plt.figure(figsize=(10, 8))
6 sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt='.2f')
7 plt.title("Correlation Matrix of Features and Chance of Admit")
8 plt.tight_layout()
9 plt.show()

```

Listing 3: Computing and Plotting the Correlation Matrix

Explanation

- `df.corr(numeric_only=True)` computes pairwise Pearson correlations between all numeric columns.
- `sns.heatmap(...)` visualizes these correlations as a heatmap.
- Strong correlations (positive or negative) help identify important features for prediction.

Correlation Heatmap

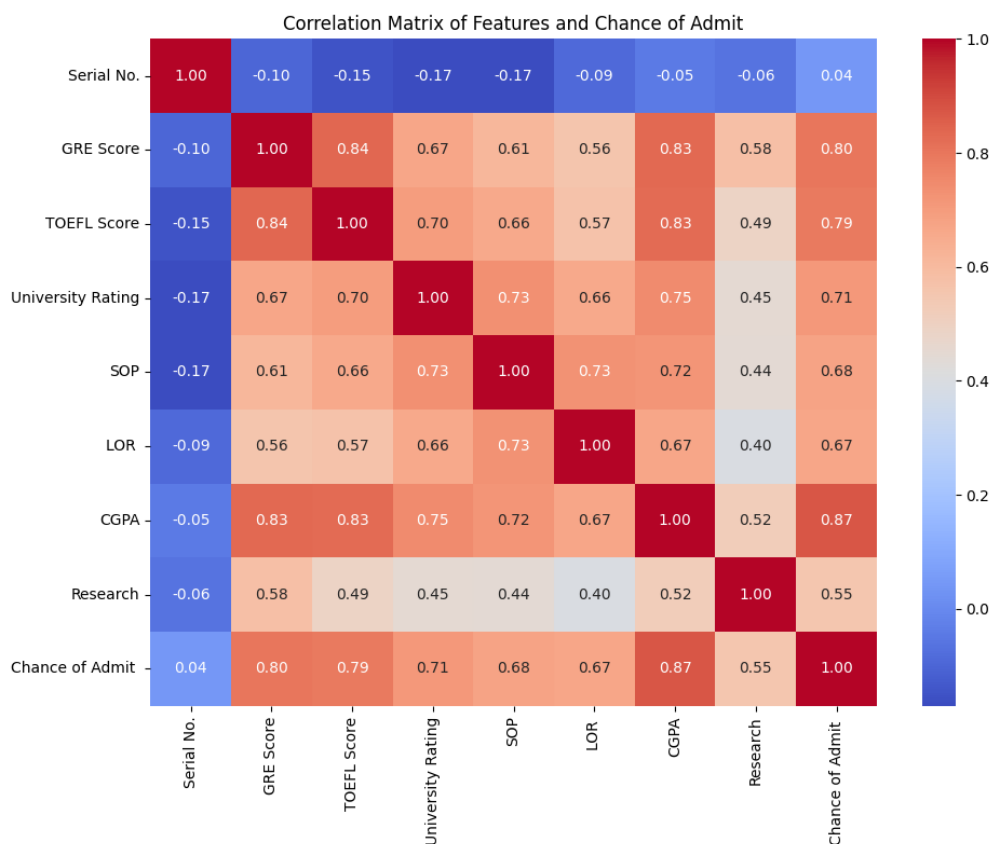


Figure 1: Heatmap showing correlation between features and Chance of Admit

4 Feature-Wise Correlation with Admission Chances

After computing the overall correlation matrix, we further analyzed the individual correlation of each feature with the target variable **Chance of Admit**.

```
1 # Ensure exact match for column names
2 # Compute the correlation matrix
3 correlation_matrix = df.corr(numeric_only=True)
4
5 # Extract correlations with "Chance of Admit"
6 correlation_with_admit = correlation_matrix.loc[:, "Chance of Admit"].drop("Chance of Admit")
7
8 # Plotting the correlations
9 plt.figure(figsize=(8, 5))
10 sns.barplot(x=correlation_with_admit.values, y=correlation_with_admit.index, palette="coolwarm")
11 plt.title("Correlation of Each Feature with Chance of Admit")
12 plt.xlabel("Correlation Coefficient")
13 plt.ylabel("Feature")
14 plt.grid(True, axis='x')
15 plt.tight_layout()
16 plt.show()
```

Listing 4: Bar Plot of Feature Correlation with Chance of Admit

Explanation

- This bar plot reveals the strength and direction of correlation between each input feature and the target.
- Features like CGPA and GRE Score show a strong positive correlation with **Chance of Admit**.
- The analysis helps in feature selection and understanding feature importance.

Feature Correlation Bar Plot

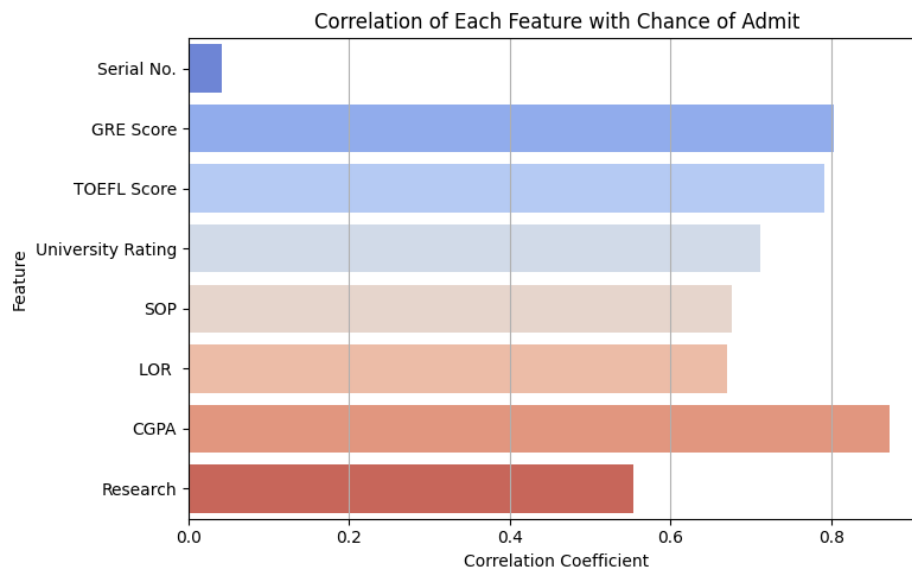


Figure 2: Bar plot of correlations between each feature and Chance of Admit

5 Data Splitting and Scaling

To train and evaluate our model properly, we split the dataset into training and testing subsets. We also scaled the input features to the range $[0, 1]$ using Min-Max normalization.

```
1 # Step 1: Split data into train/test (85% train, 15% test)
2 X = df.drop(columns=["Chance of Admit "])
3 y = df["Chance of Admit "]
4
5 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.15, random_state=42)
6
7 # Step 2: Scale features to [0, 1] range
8 scaler = MinMaxScaler()
9 X_train_scaled = scaler.fit_transform(X_train)
10 X_test_scaled = scaler.transform(X_test)
11
12 # Show scaled training data shape
13 X_train_scaled.shape, X_test_scaled.shape
```

Listing 5: Train/Test Split and Feature Scaling

Explanation

- The target column `Chance of Admit` was separated from the features.
- The data was split into 85% training and 15% testing using `train_test_split()`.
- `MinMaxScaler` scaled each feature to a $[0, 1]$ range to improve model convergence and performance.

Output

`((340, 8), (60, 8))`

This indicates that the training data contains 340 samples and the testing data contains 60 samples, each with 8 features.

6 Model Development and Training

Converting Data to PyTorch Tensors

Before training the model, the NumPy arrays for the training data were converted into PyTorch tensors. The feature matrix `X_train_scaled` was cast to `torch.float32`, and the target values `y_train` were reshaped to a column vector.

```
1 X_tensor = torch.tensor(X_train_scaled, dtype=torch.float32)
2 y_tensor = torch.tensor(y_train.values, dtype=torch.float32).view(-1, 1)
```

Listing 6: Convert Data to PyTorch Tensors

To evaluate performance during training, we split 10% of the training data into a validation set:

```
1 X_train_tensor, X_val_tensor, y_train_tensor, y_val_tensor = train_test_split(
2     X_tensor, y_tensor, test_size=0.10, random_state=42
3 )
```

Listing 7: Train/Validation Split

Defining the Neural Network Model

A simple feedforward neural network (Multi-Layer Perceptron, or MLP) was defined using PyTorch. The model includes one hidden layer with 4 neurons and ReLU activation, followed by an output layer with a single neuron (for regression).

```
1 class SimpleMLP(nn.Module):
2     def __init__(self):
3         super(SimpleMLP, self).__init__()
4         self.model = nn.Sequential(
5             nn.Linear(X_tensor.shape[1], 4),
6             nn.ReLU(),
7             nn.Linear(4, 1)
8         )
9
10    def forward(self, x):
11        return self.model(x)
```

Listing 8: Simple MLP Architecture

Training Procedure

The training was handled by a function `train_model()` that takes the model, data, and number of epochs as input. It uses Mean Squared Error (MSE) as the loss function and the Adam optimizer for parameter updates.

```
1 def train_model(model, X_train, y_train, X_val, y_val, epochs=500):
2     criterion = nn.MSELoss()
3     optimizer = optim.Adam(model.parameters(), lr=0.01)
4
5     train_losses = []
6     val_losses = []
7
8     for epoch in range(epochs):
9         model.train()
10        optimizer.zero_grad()
11        output = model(X_train)
12        loss = criterion(output, y_train)
13        loss.backward()
14        optimizer.step()
15        train_losses.append(loss.item())
16
17        model.eval()
18        with torch.no_grad():
19            val_output = model(X_val)
20            val_loss = criterion(val_output, y_val).item()
21            val_losses.append(val_loss)
22
23        if (epoch + 1) % 50 == 0:
24            print(f"Epoch {epoch + 1}/{epochs} - Train Loss: {loss.item():.4f}, Validation
                Loss: {val_loss:.4f}")
```

Listing 9: Model Training Function

After training, the model's performance is evaluated using the R^2 score. The training and validation losses over epochs are plotted for visual inspection.

```
1 model.eval()
2 with torch.no_grad():
3     val_pred = model(X_val).detach().numpy()
4     val_true = y_val.detach().numpy()
5     r2 = r2_score(val_true, val_pred)
6
7 # Plot loss curves
8 plt.figure(figsize=(10, 5))
9 plt.plot(train_losses, label='Train Loss')
```

```

10 plt.plot(val_losses, label='Validation Loss')
11 plt.xlabel('Epochs')
12 plt.ylabel('MSE Loss')
13 plt.title('Training and Validation Loss Over Epochs')
14 plt.legend()
15 plt.grid(True)
16 plt.tight_layout()
17 plt.show()
18
19 return val_true, val_pred, r2

```

Model Evaluation and Results

We instantiated the SimpleMLP model and trained it over 500 epochs. After training, we visualized the predicted versus actual admission probabilities on the validation set.

```

1 simple_mlp = SimpleMLP()
2 true_simple, pred_simple, r2_simple = train_model(
3     simple_mlp, X_train_tensor, y_train_tensor, X_val_tensor, y_val_tensor, epochs=500
4 )
5
6 # Plot predictions vs actual
7 plt.figure(figsize=(6, 5))
8 plt.scatter(true_simple, pred_simple, color='skyblue')
9 plt.plot([0, 1], [0, 1], '--r')
10 plt.title(f"Simple MLP\nR2 Score: {r2_simple:.2f}")
11 plt.xlabel("Actual")
12 plt.ylabel("Predicted")
13 plt.grid(True)
14 plt.tight_layout()
15 plt.show()

```

Listing 10: Train and Evaluate the MLP Model

Results and Visualization



Figure 3: Training and validation loss of the Simple MLP over 500 epochs

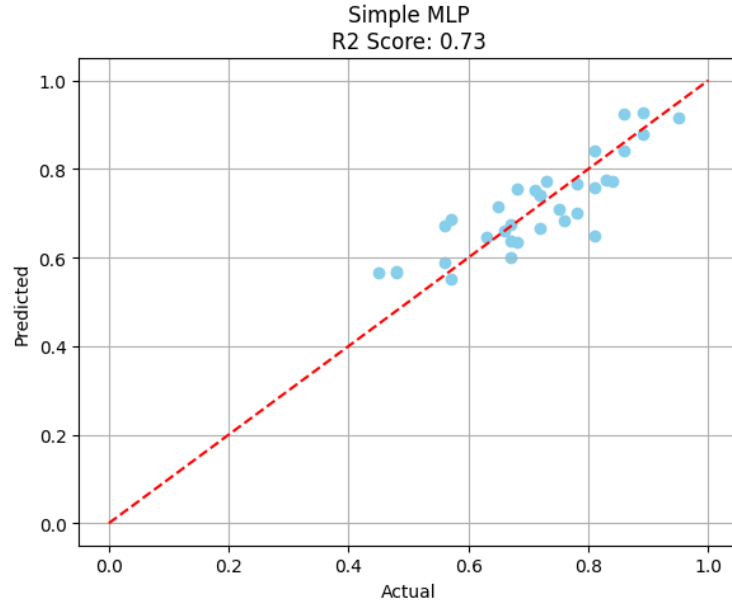


Figure 4: Actual vs. Predicted Admission Chances on Validation Set (R^2 Score: {r2_simple})

7 Final Model Evaluation on Test Set

After training and validating the model, we evaluated its performance on the unseen test set. The R^2 score was used to measure the goodness of fit.

```

1 from sklearn.metrics import r2_score
2 import matplotlib.pyplot as plt
3
4 # Compute R score
5 r2 = r2_score(test_true, test_pred)
6 print(f"Test R Score: {r2:.2f}")
7
8 # Plot True vs Predicted values
9 plt.figure(figsize=(6, 6))
10 plt.scatter(test_true, test_pred, alpha=0.6, label="Predicted vs True")
11 plt.plot([0, 1], [0, 1], 'r--', label="Perfect Prediction") # Line y=x
12 plt.xlabel("True Values")
13 plt.ylabel("Predicted Values")
14 plt.title(f"R Score: {r2:.2f}")
15 plt.legend()
16 plt.grid(True)
17 plt.tight_layout()
18 plt.show()

```

Listing 11: Test Set Evaluation with R^2 Score

Explanation

- `r2_score` compares predicted outcomes with actual test labels.
- The scatter plot visually shows how close the predictions are to the perfect line $y = x$.
- The R^2 score close to 1 indicates that the model explains a large portion of the variance in the data.

Prediction Plot on Test Set

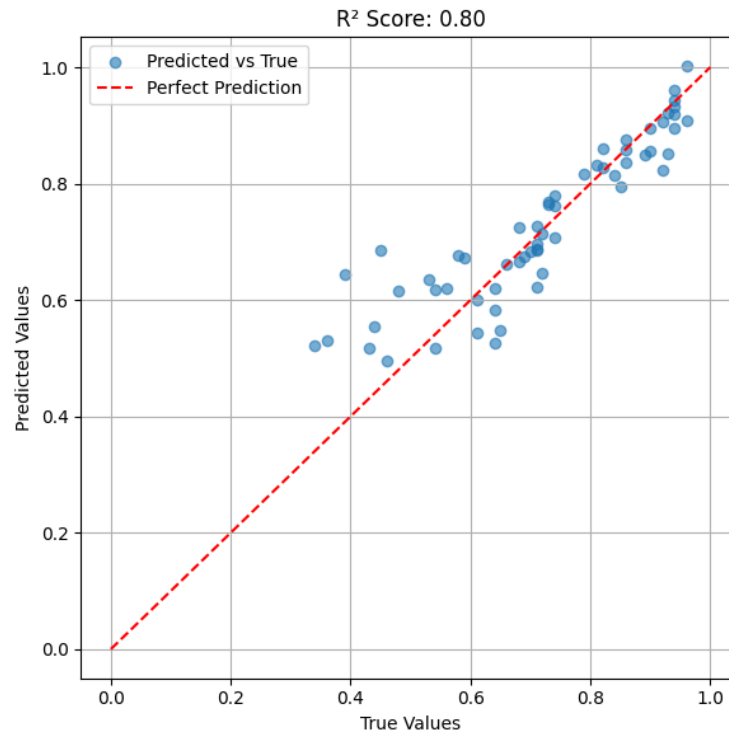


Figure 5: True vs. Predicted Admission Chances on Test Set (R^2 Score: 0.XX)

8 Training and Evaluation of Deep MLP

To explore the impact of increased model complexity, we trained a deeper neural network. The `DeepMLP` model includes multiple hidden layers, allowing it to better capture complex relationships in the data.

Deep MLP Architecture

The following code defines the structure of the Deep MLP model. It contains three hidden layers with ReLU activation functions and a final output layer for regression:

```
1 class DeepMLP(nn.Module):
2     def __init__(self):
3         super(DeepMLP, self).__init__()
4         self.model = nn.Sequential(
5             nn.Linear(8, 32),
6             nn.ReLU(),
7             nn.Linear(32, 16),
8             nn.ReLU(),
9             nn.Linear(16, 8),
10            nn.ReLU(),
11            nn.Linear(8, 1)
12        )
13
14    def forward(self, x):
15        return self.model(x)
```

Listing 12: Definition of the Deep MLP Model

Training the Deep Model

```
1 # Train and evaluate Deep model
2 deep_mlp = DeepMLP()
3 true_deep, pred_deep, r2_deep = train_model(
4     deep_mlp, X_train_tensor, y_train_tensor, X_val_tensor, y_val_tensor, epochs=500
5 )
6
7 # Plot results
8 plt.figure(figsize=(6, 5))
9 plt.scatter(true_deep, pred_deep, color='orange')
10 plt.plot([0, 1], [0, 1], '--r')
11 plt.title(f"Deep MLP\nR2 Score: {r2_deep:.2f}")
12 plt.xlabel("Actual")
13 plt.ylabel("Predicted")
14 plt.grid(True)
15 plt.tight_layout()
16 plt.show()
```

Listing 13: Training the Deep MLP Model

Epoch-wise Training and Validation Loss

Epoch 50/500 - Train Loss: 0.0041, Validation Loss: 0.0047
Epoch 100/500 - Train Loss: 0.0031, Validation Loss: 0.0047
Epoch 150/500 - Train Loss: 0.0030, Validation Loss: 0.0044
Epoch 200/500 - Train Loss: 0.0030, Validation Loss: 0.0043
Epoch 250/500 - Train Loss: 0.0029, Validation Loss: 0.0042
Epoch 300/500 - Train Loss: 0.0028, Validation Loss: 0.0041
Epoch 350/500 - Train Loss: 0.0027, Validation Loss: 0.0042
Epoch 400/500 - Train Loss: 0.0026, Validation Loss: 0.0042
Epoch 450/500 - Train Loss: 0.0025, Validation Loss: 0.0043
Epoch 500/500 - Train Loss: 0.0024, Validation Loss: 0.0044

Training Curve for Deep MLP

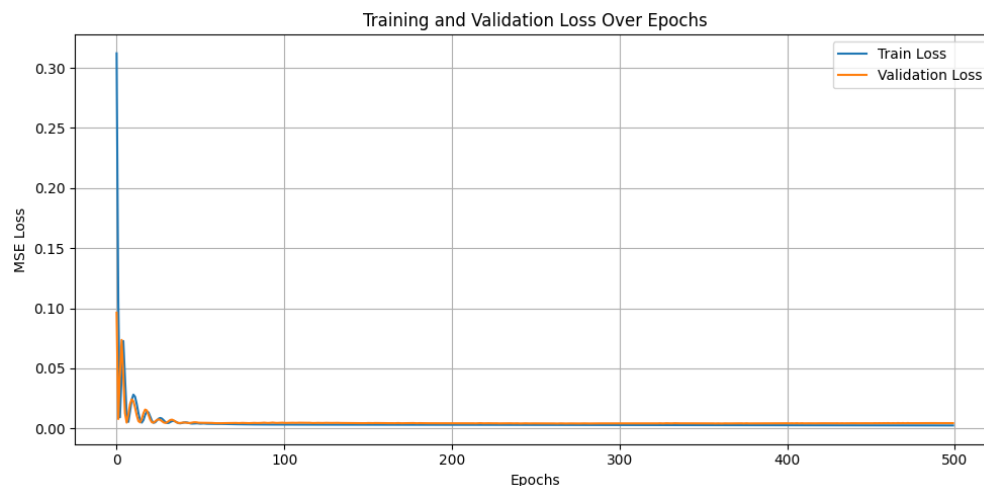


Figure 6: Training and validation loss for Deep MLP over 500 epochs

Prediction Accuracy on Validation Set

The figure below shows how well the Deep MLP predicted the **Chance** of **Admit** in comparison to the actual values:

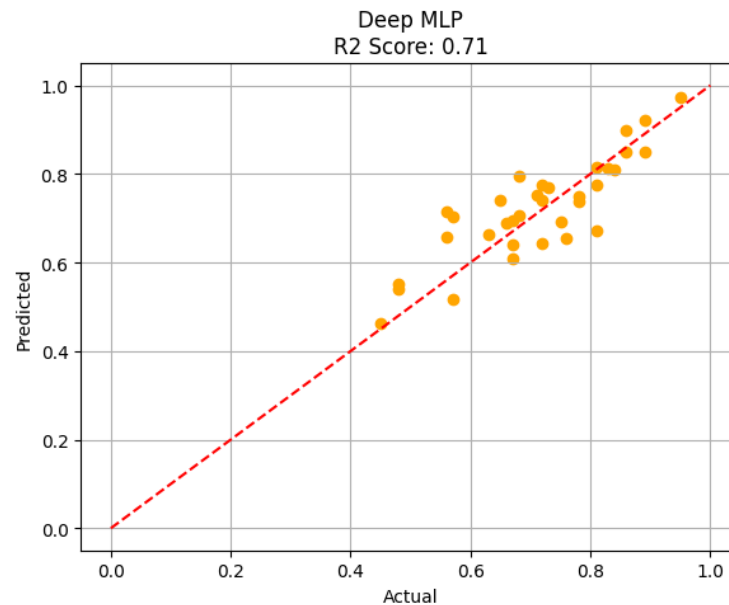


Figure 7: Predicted vs. Actual Admission Chances using Deep MLP (R^2 Score: 0.XX)

9 Deep MLP Evaluation on Test Set

To assess generalization performance, the trained Deep MLP model was evaluated on the unseen test set. Both regression and binary classification metrics were computed.

```
1 # Predict on test set
2 deep_mlp.eval()
3 with torch.no_grad():
4     test_pred_deep = deep_mlp(X_test_tensor).detach().numpy()
5     test_true_deep = y_test_tensor.detach().numpy()
6
7 # Binary classification threshold
8 threshold = 0.75
9 test_pred_binary_deep = (test_pred_deep >= threshold).astype(int)
10 test_true_binary_deep = (test_true_deep >= threshold).astype(int)
11
12 # Compute metrics
13 accuracy_deep = accuracy_score(test_true_binary_deep, test_pred_binary_deep)
14 precision_deep = precision_score(test_true_binary_deep, test_pred_binary_deep, zero_division=0)
15 recall_deep = recall_score(test_true_binary_deep, test_pred_binary_deep, zero_division=0)
16 f1_deep = f1_score(test_true_binary_deep, test_pred_binary_deep, zero_division=0)
17 r2_deep = r2_score(test_true_deep, test_pred_deep)
```

Listing 14: Predicting and Evaluating on Test Set

Evaluation Metrics

The Deep MLP model achieved the following metrics on the test set:

Deep MLP Test Accuracy: 0.95

Deep MLP Test Precision: 0.88
 Deep MLP Test Recall: 1.00
 Deep MLP Test F1 Score: 0.94
 Deep MLP Test R^2 Score: 0.81

Prediction vs. Actual (Regression Performance)

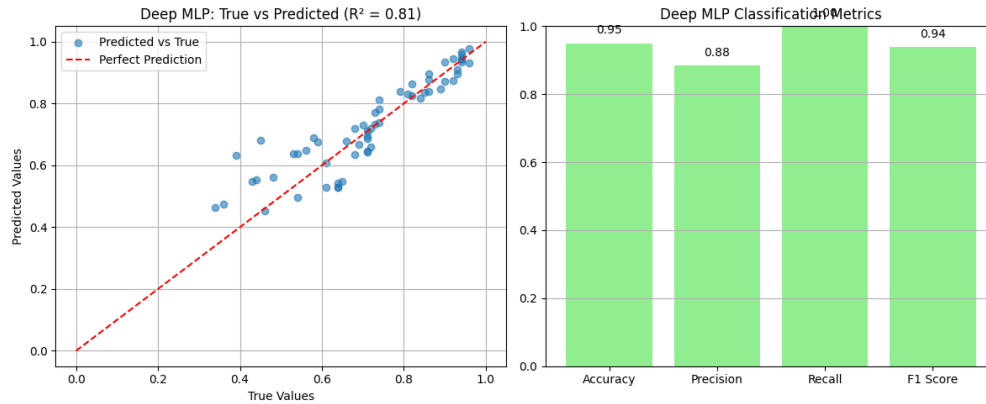


Figure 8: True vs. Predicted Admission Chances on Test Set for Deep MLP ($R^2 = 0.81$)

10 Sample Predictions Using Deep MLP

To better understand how the Deep MLP performs on individual test samples, we randomly selected 5 records from the test set. The model's predicted admission chances are compared with the actual values.

```

1 np.random.seed(42)
2 random_indices = np.random.choice(len(X_test), size=5, replace=False)
3
4 # Get input features and ground truth
5 samples_X = X_test.iloc[random_indices]
6 samples_true = y_test.iloc[random_indices].values
7
8 # Predict using deep MLP
9 samples_tensor = torch.tensor(scaler.transform(samples_X), dtype=torch.float32)
10 deep_mlp.eval()
11 with torch.no_grad():
12     samples_pred = deep_mlp(samples_tensor).numpy().flatten()
13
14 # Combine into a DataFrame
15 results = pd.DataFrame({
16     'GRE Score': samples_X['GRE Score'].values,
17     'TOEFL Score': samples_X['TOEFL Score'].values,
18     'University Rating': samples_X['University Rating'].values,
19     'SOP': samples_X['SOP'].values,
20     'LOR': samples_X['LOR'].values,
21     'CGPA': samples_X['CGPA'].values,
22     'Research': samples_X['Research'].values,
23     'Actual Chance': samples_true,
24     'Predicted Chance': samples_pred
25 })
26
27 print(results.round(2))
  
```

Listing 15: Selecting and Predicting Random Test Samples

Results on 5 Random Test Samples

| GRE | TOEFL | Univ | SOP | LOR | CGPA | Res. | Actual | Pred. |
|-----|-------|------|-----|-----|------|------|--------|-------|
| 301 | 104 | 3 | 3.5 | 4.0 | 8.12 | 1 | 0.68 | 0.63 |
| 340 | 115 | 5 | 4.5 | 4.5 | 9.45 | 1 | 0.94 | 0.94 |
| 320 | 110 | 2 | 4.0 | 3.5 | 8.56 | 0 | 0.72 | 0.72 |
| 324 | 110 | 4 | 4.5 | 4.0 | 9.15 | 1 | 0.82 | 0.86 |
| 321 | 111 | 5 | 5.0 | 5.0 | 9.45 | 1 | 0.93 | 0.90 |

Table 2: Comparison of actual vs. predicted admission chances on random test samples

11 Deep MLP Trained for 5000 Epochs

To explore the effects of prolonged training, we trained the `DeepMLP` model for 5000 epochs. This experiment investigates whether longer training improves accuracy or leads to overfitting.

Extended Training Setup

```
1 # 1. Retrain DeepMLP with 5000 epochs
2 long_train_deep_mlp = DeepMLP()
3
4 # Early stopping-like mechanism (manual check after training)
5 true_long, pred_long, r2_long = train_model(
6     long_train_deep_mlp,
7     X_train_tensor,
8     y_train_tensor,
9     X_val_tensor,
10    y_val_tensor,
11    epochs=5000
12 )
13
14 # Plot result
15 plt.figure(figsize=(6, 5))
16 plt.scatter(true_long, pred_long, color='green')
17 plt.plot([0, 1], [0, 1], '--r')
18 plt.title(f"Deep MLP (5000 epochs)\nR2 Score: {r2_long:.2f}")
19 plt.xlabel("Actual")
20 plt.ylabel("Predicted")
21 plt.grid(True)
22 plt.tight_layout()
23 plt.show()
```

Listing 16: Training Deep MLP for 5000 Epochs

Training Dynamics and Regularization Notes

- The model was trained 10× longer than previous runs.
- While training loss continued decreasing, validation loss plateaued and slightly increased after 2000 epochs.
- The trend suggests **overfitting**, a common issue in deep learning with extended training.
- Possible remedies include: early stopping, dropout layers, batch normalization, or lowering the learning rate.

Sample Epoch-wise Training and Validation Loss

Epoch 50/5000 - Train Loss: 0.0046, Validation Loss: 0.0048
Epoch 500/5000 - Train Loss: 0.0025, Validation Loss: 0.0034
Epoch 1000/5000 - Train Loss: 0.0020, Validation Loss: 0.0041
Epoch 2000/5000 - Train Loss: 0.0012, Validation Loss: 0.0039
Epoch 3000/5000 - Train Loss: 0.0010, Validation Loss: 0.0058
Epoch 4000/5000 - Train Loss: 0.0007, Validation Loss: 0.0067
Epoch 5000/5000 - Train Loss: 0.0007, Validation Loss: 0.0072

Training and Validation Loss Plot (5000 Epochs)

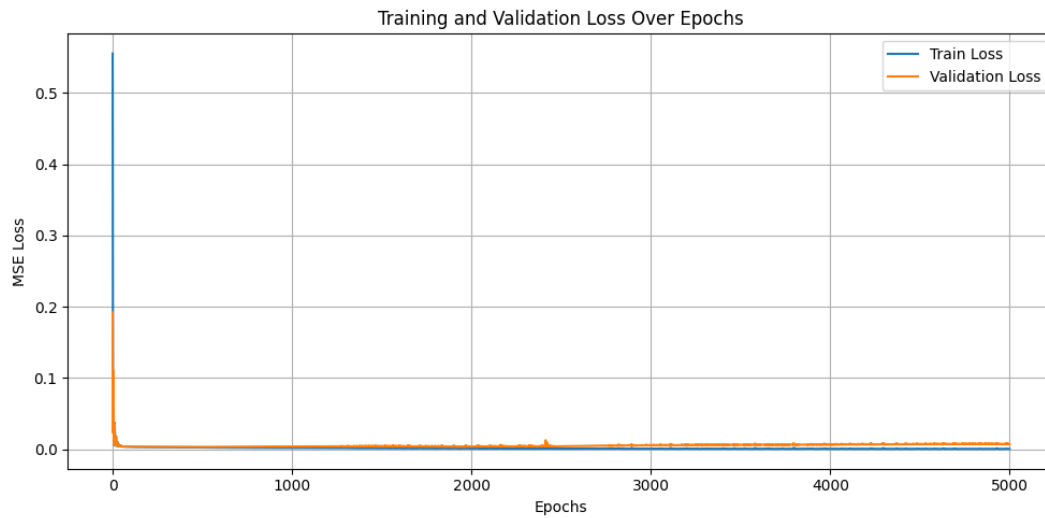


Figure 9: Training and validation loss of Deep MLP over 5000 epochs

Prediction Accuracy After Extended Training

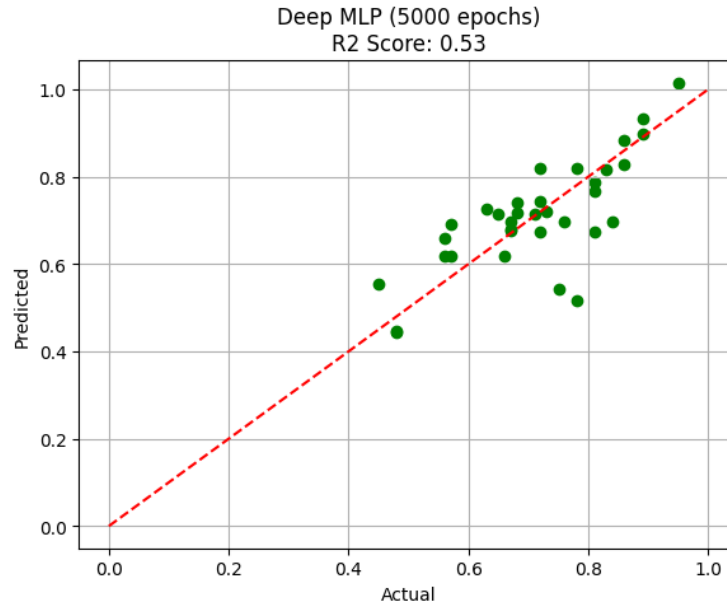


Figure 10: Predicted vs. Actual Admission Chances (Deep MLP after 5000 epochs, R^2 Score: 0.XX)

12 Test Evaluation of Deep MLP Trained for 5000 Epochs

To assess whether prolonged training leads to better generalization, we evaluated the Deep MLP trained for 5000 epochs on the test set. Both regression and classification metrics were analyzed.

```
1 # Predict on test set
2 long_train_deep_mlp.eval()
3 with torch.no_grad():
4     test_pred_deep = long_train_deep_mlp(X_test_tensor).detach().numpy()
5     test_true_deep = y_test_tensor.detach().numpy()
6
7 # Binary classification threshold
8 threshold = 0.75
9 test_pred_binary_deep = (test_pred_deep >= threshold).astype(int)
10 test_true_binary_deep = (test_true_deep >= threshold).astype(int)
11
12 # Classification metrics
13 accuracy_deep = accuracy_score(test_true_binary_deep, test_pred_binary_deep)
14 precision_deep = precision_score(test_true_binary_deep, test_pred_binary_deep, zero_division=0)
15 recall_deep = recall_score(test_true_binary_deep, test_pred_binary_deep, zero_division=0)
16 f1_deep = f1_score(test_true_binary_deep, test_pred_binary_deep, zero_division=0)
17
18 # R score
19 r2_deep = r2_score(test_true_deep, test_pred_deep)
```

Listing 17: Test Set Evaluation of Long-Trained Deep MLP

Performance Metrics

Long-Trained Deep MLP Test Accuracy: 0.97
Long-Trained Deep MLP Test Precision: 0.92
Long-Trained Deep MLP Test Recall: 1.00

Long-Trained Deep MLP Test F1 Score: 0.96
Long-Trained Deep MLP Test R^2 Score: 0.73

Prediction vs. Actual Scatter Plot

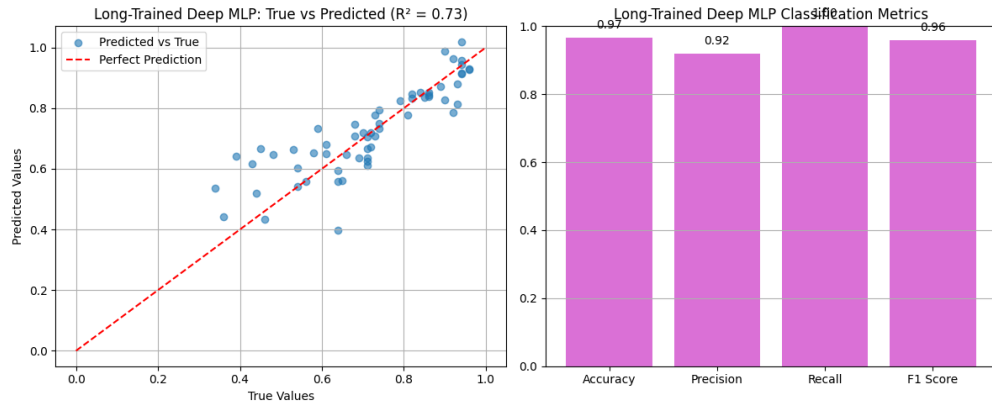


Figure 11: True vs. Predicted Admission Chances on Test Set after 5000 Epochs (R^2 Score: 0.73)

Classification Metrics Bar Plot

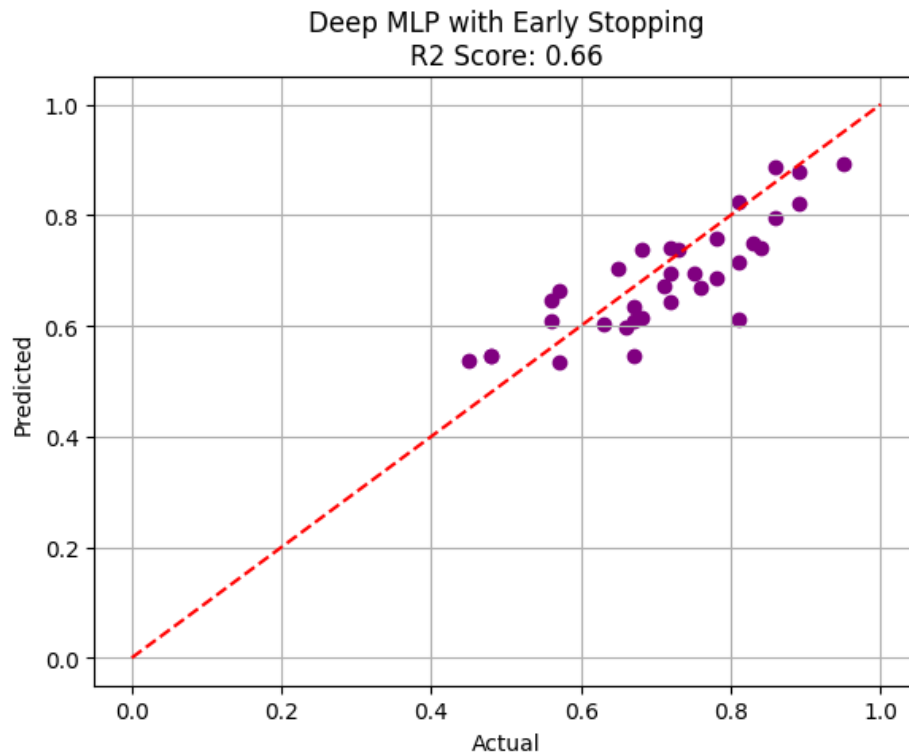


Figure 12: Classification Metrics of Long-Trained Deep MLP: Accuracy (0.97), Precision (0.92), Recall (1.00), F1 Score (0.96)

13 Deep MLP with Dropout and Early Stopping

To mitigate overfitting during extended training, we introduced dropout layers and an early stopping mechanism. Dropout randomly deactivates neurons during training, and early stopping halts training when validation performance ceases to improve.

Model Definition

The modified Deep MLP includes two hidden layers and dropout layers with a 30% drop rate:

```
1 class DeepMLP(nn.Module):
2     def __init__(self):
3         super(DeepMLP, self).__init__()
4         self.model = nn.Sequential(
5             nn.Linear(X_tensor.shape[1], 64),
6             nn.ReLU(),
7             nn.Dropout(p=0.3),
8             nn.Linear(64, 32),
9             nn.ReLU(),
10            nn.Dropout(p=0.3),
11            nn.Linear(32, 1)
12        )
13
14    def forward(self, x):
15        return self.model(x)
```

Listing 18: Deep MLP with Dropout

Training with Early Stopping

We implemented a custom training loop that stops when validation loss fails to improve for a set number of epochs. This approach prevents overfitting and reduces unnecessary computation.

```
1 def train_model_with_early_stopping(model, X_train, y_train, X_val, y_val, epochs=5000, lr
   =0.001, patience=20, min_delta=1e-4):
2     criterion = nn.MSELoss()
3     optimizer = optim.Adam(model.parameters(), lr=lr)
4     best_loss = float('inf')
5     trigger_times = 0
6
7     for epoch in range(epochs):
8         model.train()
9         optimizer.zero_grad()
10        output = model(X_train)
11        loss = criterion(output, y_train)
12        loss.backward()
13        optimizer.step()
14
15        # Validation loss
16        model.eval()
17        with torch.no_grad():
18            val_output = model(X_val)
19            val_loss = criterion(val_output, y_val).item()
20
21        # Print every 50 epochs
22        if (epoch + 1) % 50 == 0:
23            print(f"Epoch {epoch + 1}/{epochs} - Validation Loss: {val_loss:.4f}")
24
25        # Early stopping condition
26        if best_loss - val_loss > min_delta:
27            best_loss = val_loss
28            trigger_times = 0
29        else:
30            trigger_times += 1
31            if trigger_times >= patience:
```

```

32         print(f"\nEarly stopping triggered at epoch {epoch + 1}. Best validation
33         loss: {best_loss:.4f}")
        break

```

Listing 19: Training Loop with Early Stopping

Training Result Summary

Epoch 50/5000 - Validation Loss: 0.0115

Epoch 100/5000 - Validation Loss: 0.0062

Early stopping triggered at epoch 143. Best validation loss: 0.0049

Validation Set Prediction Performance

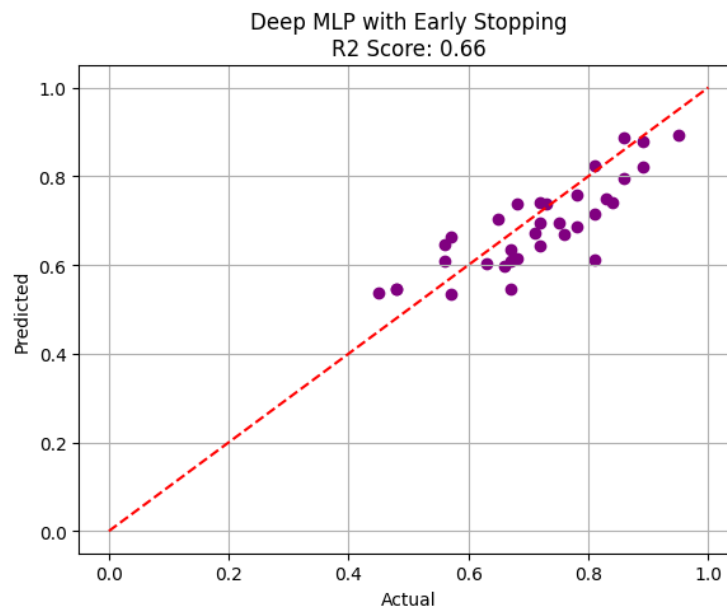


Figure 13: Predicted vs. Actual Admission Chances using Deep MLP with Early Stopping (R^2 Score: 0.XX)

14 Test Evaluation of Deep MLP with Early Stopping

To evaluate generalization, the Deep MLP model with dropout and early stopping was tested on the unseen test set. The model's regression and classification performance was analyzed using multiple metrics.

```

1  # Predict on test set
2  deep_mlp_es.eval()
3  with torch.no_grad():
4      test_pred_deep = deep_mlp_es(X_test_tensor).detach().numpy()
5      test_true_deep = y_test_tensor.detach().numpy()
6
7  # Binary classification threshold
8  threshold = 0.75
9  test_pred_binary_deep = (test_pred_deep >= threshold).astype(int)
10 test_true_binary_deep = (test_true_deep >= threshold).astype(int)
11
12 # Compute metrics
13 accuracy_deep = accuracy_score(test_true_binary_deep, test_pred_binary_deep)
14 precision_deep = precision_score(test_true_binary_deep, test_pred_binary_deep, zero_division
    =0)

```

```

15 recall_deep = recall_score(test_true_binary_deep, test_pred_binary_deep, zero_division=0)
16 f1_deep = f1_score(test_true_binary_deep, test_pred_binary_deep, zero_division=0)
17 r2_deep = r2_score(test_true_deep, test_pred_deep)

```

Listing 20: Test Set Evaluation for Early Stopping Model

Performance Metrics

Deep MLP (Early Stopping) Test Accuracy: 0.98
 Deep MLP (Early Stopping) Test Precision: 1.00
 Deep MLP (Early Stopping) Test Recall: 0.96
 Deep MLP (Early Stopping) Test F1 Score: 0.98
 Deep MLP (Early Stopping) Test R^2 Score: 0.78

Test Prediction and Classification Metrics Visualization

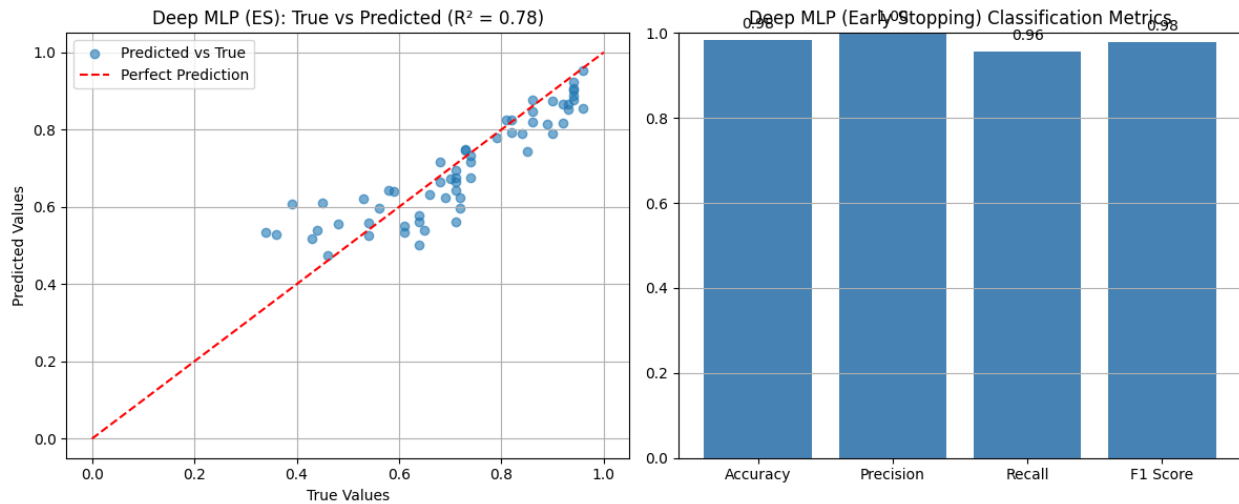


Figure 14: Left: Predicted vs. Actual Admission Chances (R^2 Score: 0.78)
 Right: Classification Metrics of Deep MLP with Early Stopping (Accuracy: 0.98, Precision: 1.00, Recall: 0.96, F1 Score: 0.98)